

# Report on Function Approximation with Neural Network and Backpropagation

1905113 - Anamul Hoque Emtiaj

## 1 Setup and Execution Instructions

### 1.1 Dependencies Installation

To run the project, make sure to have Python installed and install the following dependencies:

```
pip install numpy pandas matplotlib tqdm torchvision scikit-learn
```

This will install essential packages required for handling datasets like FashionMNIST.

### 1.2 Training the Model

To train the model, follow these steps:

- Place the 1905113.ipynb file in your project folder.
- Specify the model architecture using the following configuration:

```
config = {
    'learning_rate': 0.0005,
    'num_epochs': 50,
    'batch_size': 32,
    'hidden_layers': [
        {'size': 512, 'bn': True, 'dropout': 0.3, 'relu': True},
        {'size': 256, 'bn': False, 'dropout': 0.4, 'relu': True},
        {'size': 128, 'bn': True, 'dropout': 0.5, 'relu': True}
    ]
}
```

- Build the model by calling the function `build_model(config)`, which will return the model.
- Train the model using the following command:

```
model.fit(train_data, train_labels, val_data, val_labels)
```

### 1.3 Testing with Saved Model

To test the model with pre-trained parameters, proceed with the following steps:

- Place the `model_1905113.pickle` file in the project folder.
- Open the 1905113.ipynb file, comment out the 'Validation', 'Train and Save Best model' and 'Independent Testing - Direct Trained Model' sections.
- Load the previously saved model using the following command:

```
loaded_model = load_model('model_1905113.pkl')
```

- Make predictions on the test data with the following commands:

```

test_predictions = loaded_model.predict(X_test)
test_loss, test_acc, test_macro_f1 = loaded_model.evaluate(X_test, y_test_one_hot)
print(f"Test Loss: {test_loss:.4f}, Test Accuracy: {test_acc:.4f}, Test Macro F1: {test_macro_f1:.4f}")

conf_matrix = confusion_matrix(y_test, test_predictions)
print("Confusion Matrix:")
print(conf_matrix)

```

## 2 Validation Performance

### 2.1 With Three Different Architectures

For this section, we experimented with three different neural network architectures. The configuration of each model and their performance results are described below.

#### 2.1.1 Model 1

##### Configuration:

```

config_model_1 = {
    'learning_rate': 0.001,
    'num_epochs': 50,
    'batch_size': 32,
    'hidden_layers': [
        {'size': 512, 'bn': True, 'dropout': 0.3, 'relu': True},
        {'size': 256, 'bn': False, 'dropout': 0.4, 'relu': True},
        {'size': 128, 'bn': True, 'dropout': 0.5, 'relu': True}
    ]
}

```

##### Performance:

- Train Accuracy: 0.9366
- Validation Accuracy: 0.9028
- Train Loss: 0.3159
- Validation Loss: 0.40
- Validation Macro F1 Score: 0.9039

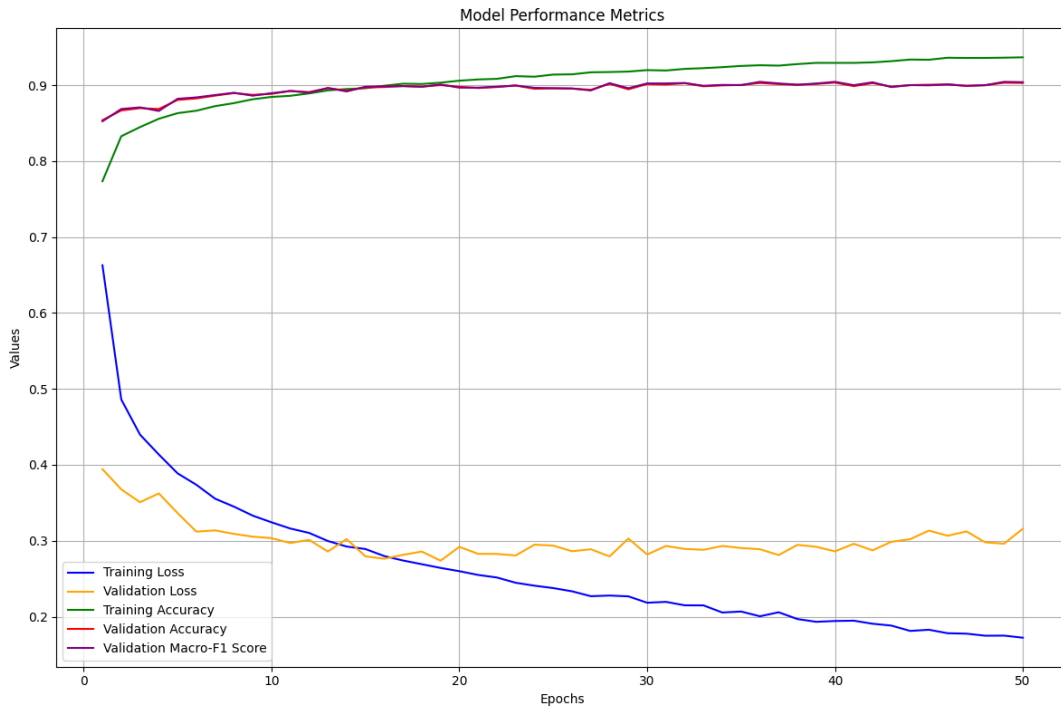


Figure 1: Training and Validation History for Model 1

### Confusion Matrix:

Table 1: Confusion Matrix

	0	1	2	3	4	5	6	7	8	9
0	519	2	13	9	3	0	50	0	1	0
1	1	597	1	8	1	0	0	0	0	0
2	10	0	524	2	49	0	26	0	0	0
3	31	3	6	513	23	0	8	0	3	0
4	1	0	36	12	548	0	27	0	3	0
5	0	0	0	0	0	600	0	16	1	4
6	70	1	47	9	33	0	455	0	4	0
7	0	0	0	0	0	7	0	536	0	7
8	2	0	3	1	3	0	12	1	568	0
9	0	0	0	0	0	8	0	25	0	557

### 2.1.2 Model 2

#### Configuration:

```
config_model_2 = {
    'learning_rate': 0.001,
    'num_epochs': 50,
    'batch_size': 64,
    'hidden_layers': [
        {'size': 512, 'bn': False, 'dropout': 0.3, 'relu': True},
        {'size': 256, 'bn': True, 'dropout': 0.4, 'relu': True}
    ]
}
```

#### Performance:

- Train Accuracy: 0.9336

- Validation Accuracy: 0.9025
- Train Loss: 0.1760
- Validation Loss: 0.2910
- Validation Macro F1 Score: 0.9033

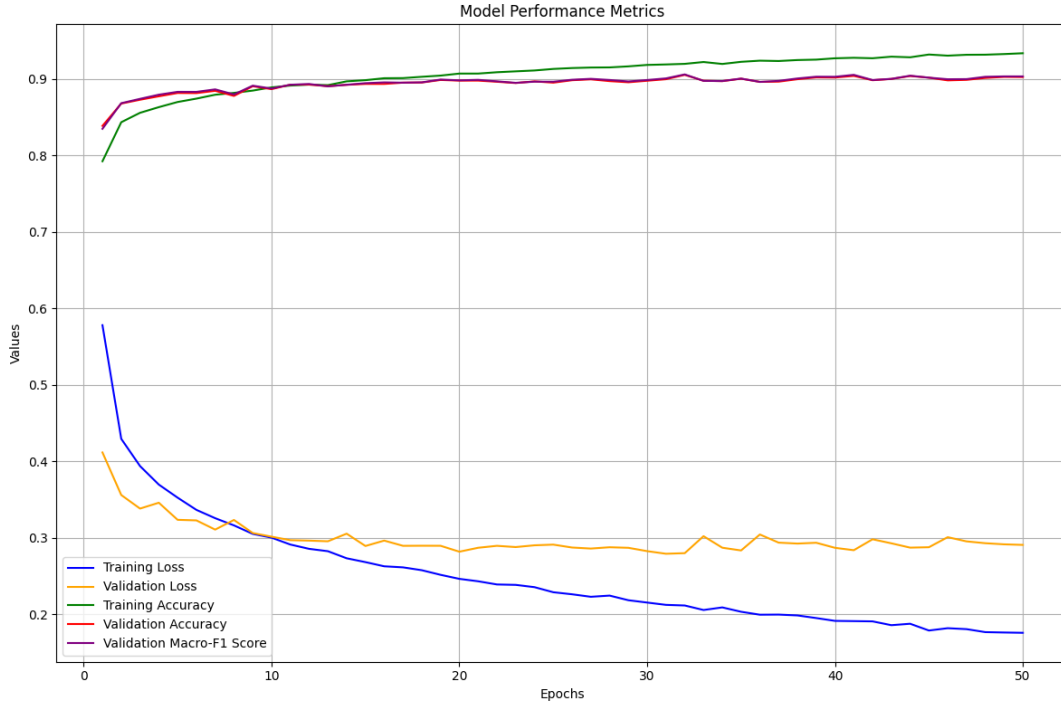


Figure 2: Training and Validation History for Model 2

### Confusion Matrix:

Table 2: Confusion Matrix

	0	1	2	3	4	5	6	7	8	9
0	525	0	5	16	2	0	47	0	2	0
1	1	595	0	12	0	0	0	0	0	0
2	16	0	500	4	57	0	34	0	0	0
3	13	2	2	555	13	0	2	0	0	0
4	2	0	37	27	521	0	38	0	2	0
5	0	0	0	0	0	602	0	13	0	6
6	75	1	36	17	31	0	454	0	5	0
7	0	0	0	0	0	6	0	529	0	15
8	4	0	3	2	2	1	6	1	571	0
9	0	0	0	0	0	12	0	15	0	563

### 2.1.3 Model 3

#### Configuration:

```
config_model_3 = {
    'learning_rate': 0.001,
    'num_epochs': 50,
    'batch_size': 128,
    'hidden_layers': [
```

```

    {'size': 1024, 'bn': True, 'dropout': 0.3, 'relu': True},
    {'size': 512, 'bn': True, 'dropout': 0.5, 'relu': True}
]
}

```

#### Performance:

- Train Accuracy: 0.9672
- Validation Accuracy: 0.9015
- Train Loss: 0.0856
- Validation Loss: 0.3663
- Validation Macro F1 Score: 0.9021

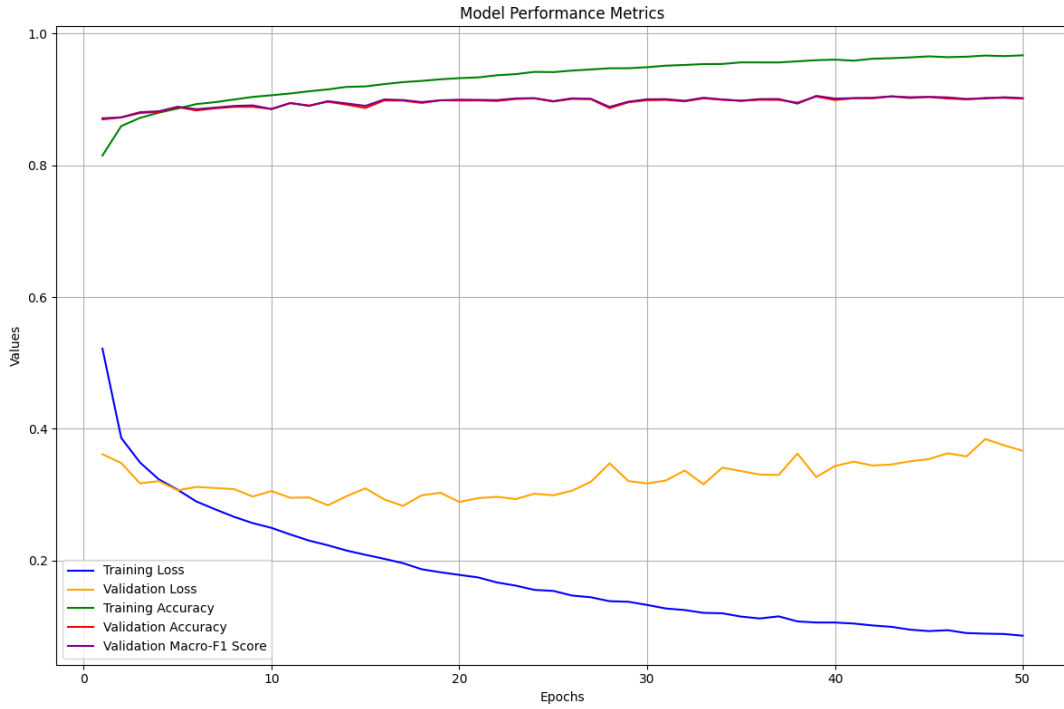


Figure 3: Training and Validation History for Model 3

#### Confusion Matrix:

Table 3: Confusion Matrix

	0	1	2	3	4	5	6	7	8	9
0	553	0	11	3	2	0	25	0	3	0
1	1	595	1	11	0	0	0	0	0	0
2	13	0	535	2	32	0	29	0	0	0
3	29	0	5	528	14	0	9	1	1	0
4	3	0	67	18	498	0	36	0	5	0
5	0	0	0	0	0	595	0	14	1	11
6	94	1	58	16	12	0	429	0	9	0
7	0	0	0	0	0	6	0	536	0	8
8	2	1	2	1	2	0	4	2	576	0
9	0	0	0	0	0	6	0	20	0	564

### 2.1.4 Comparison of Architectures

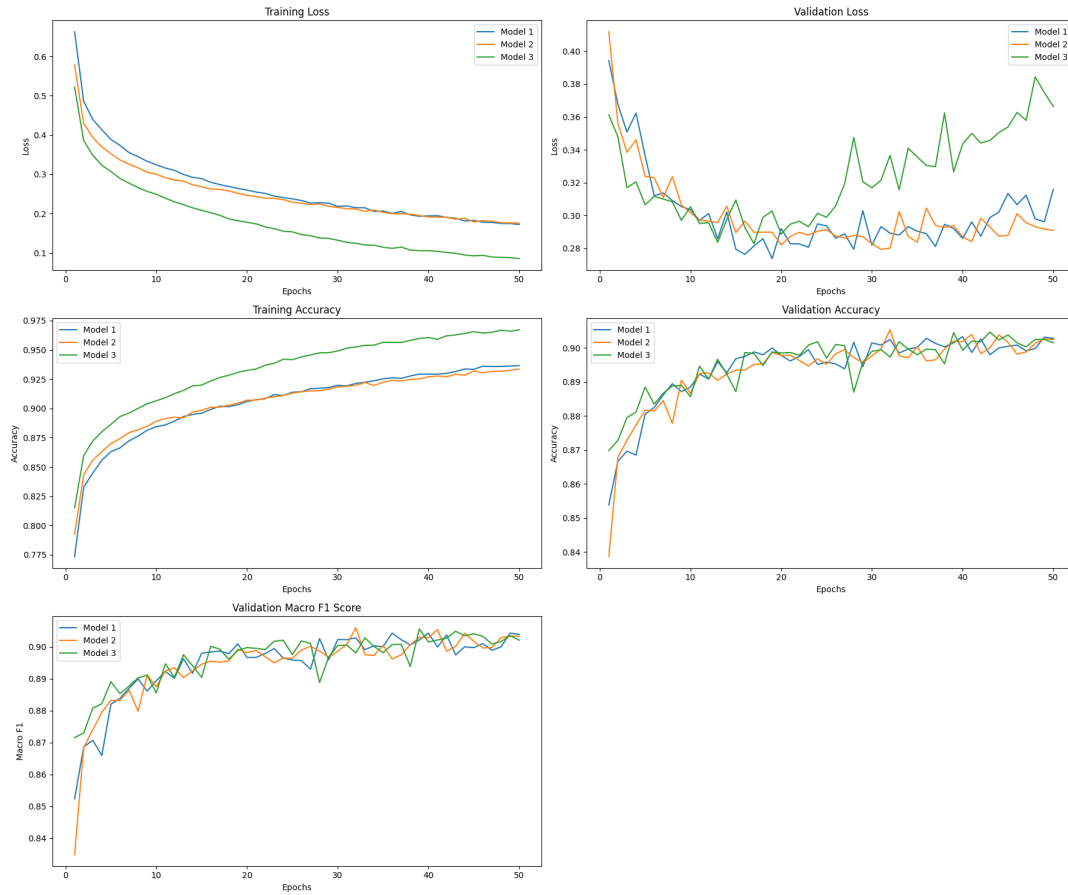


Figure 4: Comparison of Training and Validation History for All Three Architectures

## 2.2 With Four Different Learning Rates

We experimented with four different learning rates. Below are the results.

### 2.2.1 Model 1 (Learning Rate = 0.005)

**Configuration:**

```
config_model_1 = {
    'learning_rate': 0.005,
    'num_epochs': 50,
    'batch_size': 32,
    'hidden_layers': [
        {'size': 512, 'bn': True, 'dropout': 0.3, 'relu': True},
        {'size': 256, 'bn': False, 'dropout': 0.4, 'relu': True},
        {'size': 128, 'bn': True, 'dropout': 0.5, 'relu': True}
    ]
}
```

**Performance:**

- Train Accuracy: 0.9281
- Validation Accuracy: 0.8992
- Train Loss: 0.2005

- Validation Loss: 0.2995
- Validation Macro F1 Score: 0.8992

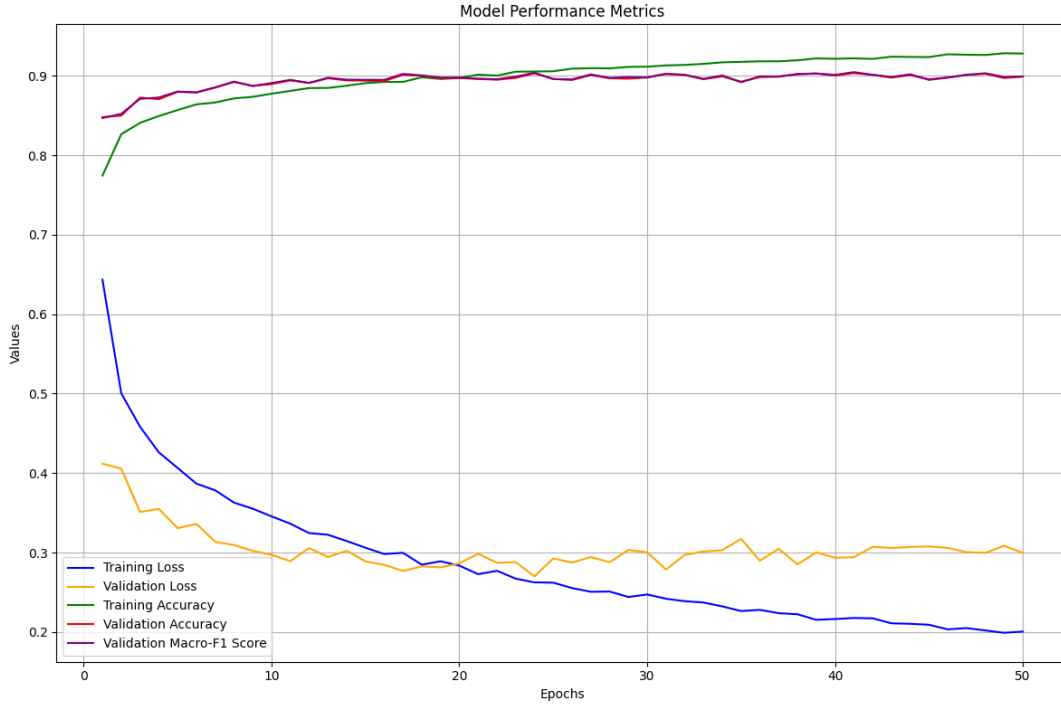


Figure 5: Training and Validation History for Model 1 (Learning Rate = 0.005)

### Confusion Matrix:

Table 4: Confusion Matrix

	0	1	2	3	4	5	6	7	8	9
0	542	1	7	17	1	0	25	0	4	0
1	1	600	0	7	0	0	0	0	0	0
2	17	0	539	4	26	0	25	0	0	0
3	15	2	2	533	24	0	8	0	3	0
4	2	0	60	12	514	0	37	0	2	0
5	0	0	0	0	0	603	0	10	1	7
6	108	1	63	13	28	0	399	0	7	0
7	0	0	0	0	0	7	0	535	0	8
8	4	1	3	2	2	0	7	1	570	0
9	0	0	0	0	0	7	0	23	0	560

### 2.2.2 Model 2 (Learning Rate = 0.0001)

#### Configuration:

```
config_model_2 = {
    'learning_rate': 0.0001,
    'num_epochs': 50,
    'batch_size': 32,
    'hidden_layers': [
        {'size': 512, 'bn': True, 'dropout': 0.3, 'relu': True},
        {'size': 256, 'bn': False, 'dropout': 0.4, 'relu': True},
        {'size': 128, 'bn': True, 'dropout': 0.5, 'relu': True}
    ]
}
```

```
]
}
```

### Performance:

- Train Accuracy: 0.9267
- Validation Accuracy: 0.9003
- Train Loss: 0.2025
- Validation Loss: 0.2774
- Validation Macro F1 Score: 0.9010

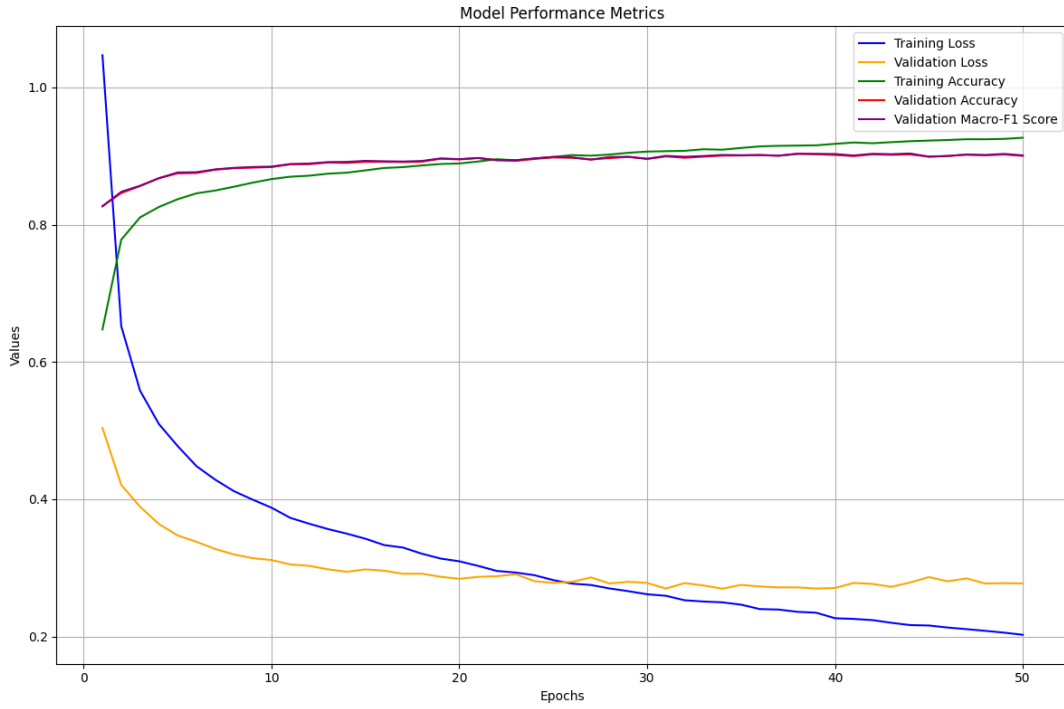


Figure 6: Training and Validation History for Model 2 (Learning Rate = 0.0001)

### Confusion Matrix:

Table 5: Confusion Matrix

	0	1	2	3	4	5	6	7	8	9
0	524	1	7	18	1	0	45	0	1	0
1	1	597	1	9	0	0	0	0	0	0
2	18	0	521	5	35	0	32	0	0	0
3	16	0	2	545	18	0	5	0	1	0
4	0	0	47	16	523	0	39	0	2	0
5	0	0	0	0	0	605	0	10	1	5
6	81	2	56	17	26	0	432	0	5	0
7	0	0	0	0	0	10	0	530	0	10
8	2	0	2	2	3	1	8	2	570	0
9	0	0	0	0	0	10	0	25	0	555



### 2.2.3 Model 3 (Learning Rate = 0.0005)

#### Configuration:

```
config_model_3 = {  
    'learning_rate': 0.0005,  
    'num_epochs': 50,  
    'batch_size': 32,  
    'hidden_layers': [  
        {'size': 512, 'bn': True, 'dropout': 0.3, 'relu': True},  
        {'size': 256, 'bn': False, 'dropout': 0.4, 'relu': True},  
        {'size': 128, 'bn': True, 'dropout': 0.5, 'relu': True}  
    ]  
}
```

#### Performance:

- Train Accuracy: 0.9367
- Validation Accuracy: 0.9035
- Train Loss: 0.1724
- Validation Loss: 0.2959
- Validation Macro F1 Score: 0.9034

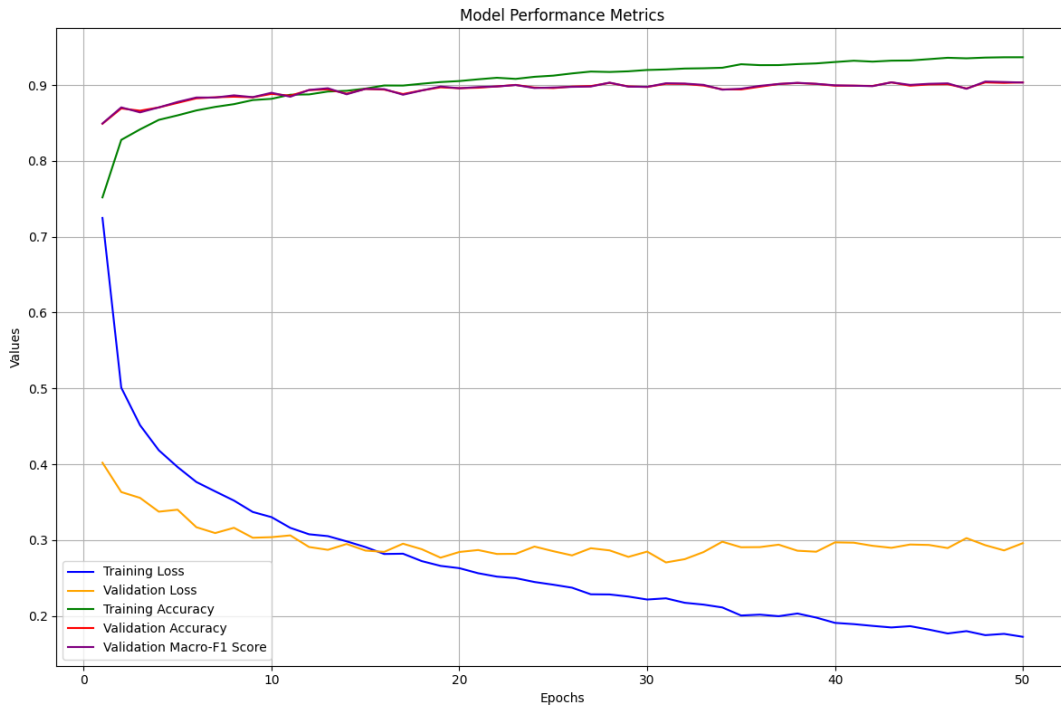


Figure 7: Training and Validation History for Model 3 (Learning Rate = 0.0005)

#### Confusion Matrix:

Table 6: Confusion Matrix

	0	1	2	3	4	5	6	7	8	9
0	547	0	7	11	1	0	27	0	4	0
1	1	599	0	8	0	0	0	0	0	0
2	18	0	521	5	47	0	20	0	0	0
3	29	2	1	537	12	0	4	0	2	0
4	1	0	38	22	542	0	20	0	4	0
5	0	0	0	0	0	598	0	12	2	9
6	102	3	45	16	39	0	408	0	6	0
7	0	0	0	0	0	9	0	535	0	6
8	1	1	0	3	2	0	6	1	576	0
9	0	0	0	0	0	7	0	25	0	558

#### 2.2.4 Model 4 (Learning Rate = 0.00001)

##### Configuration:

```
config_model_4 = {
    'learning_rate': 0.00001,
    'num_epochs': 50,
    'batch_size': 32,
    'hidden_layers': [
        {'size': 512, 'bn': True, 'dropout': 0.3, 'relu': True},
        {'size': 256, 'bn': False, 'dropout': 0.4, 'relu': True},
        {'size': 128, 'bn': True, 'dropout': 0.5, 'relu': True}
    ]
}
```

##### Performance:

- Train Accuracy: 0.8512
- Validation Accuracy: 0.8778
- Train Loss: 0.4281
- Validation Loss: 0.3324
- Validation Macro F1 Score: 0.8791

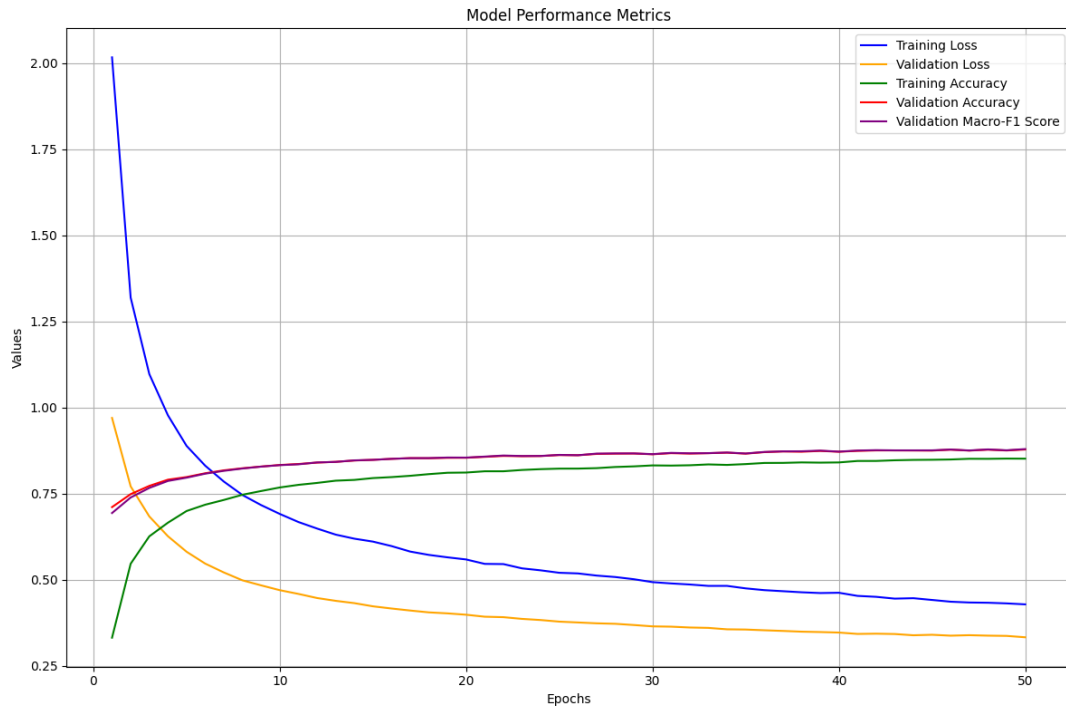


Figure 8: Training and Validation History for Model 4 (Learning Rate = 0.00001)

### Confusion Matrix:

Table 7: Confusion Matrix										
	0	1	2	3	4	5	6	7	8	9
0	503	0	10	20	1	0	62	0	1	0
1	1	587	1	16	2	0	1	0	0	0
2	5	0	497	3	59	0	47	0	0	0
3	21	2	4	525	23	0	9	0	2	1
4	0	0	45	21	518	0	39	0	4	0
5	0	0	0	0	0	595	0	16	1	9
6	77	0	63	13	47	0	411	0	8	0
7	0	0	0	0	0	19	0	514	1	16
8	1	1	3	2	3	2	16	3	559	0
9	0	0	0	0	0	10	0	22	0	558

### 2.2.5 Comparison of Models with different Learning rate

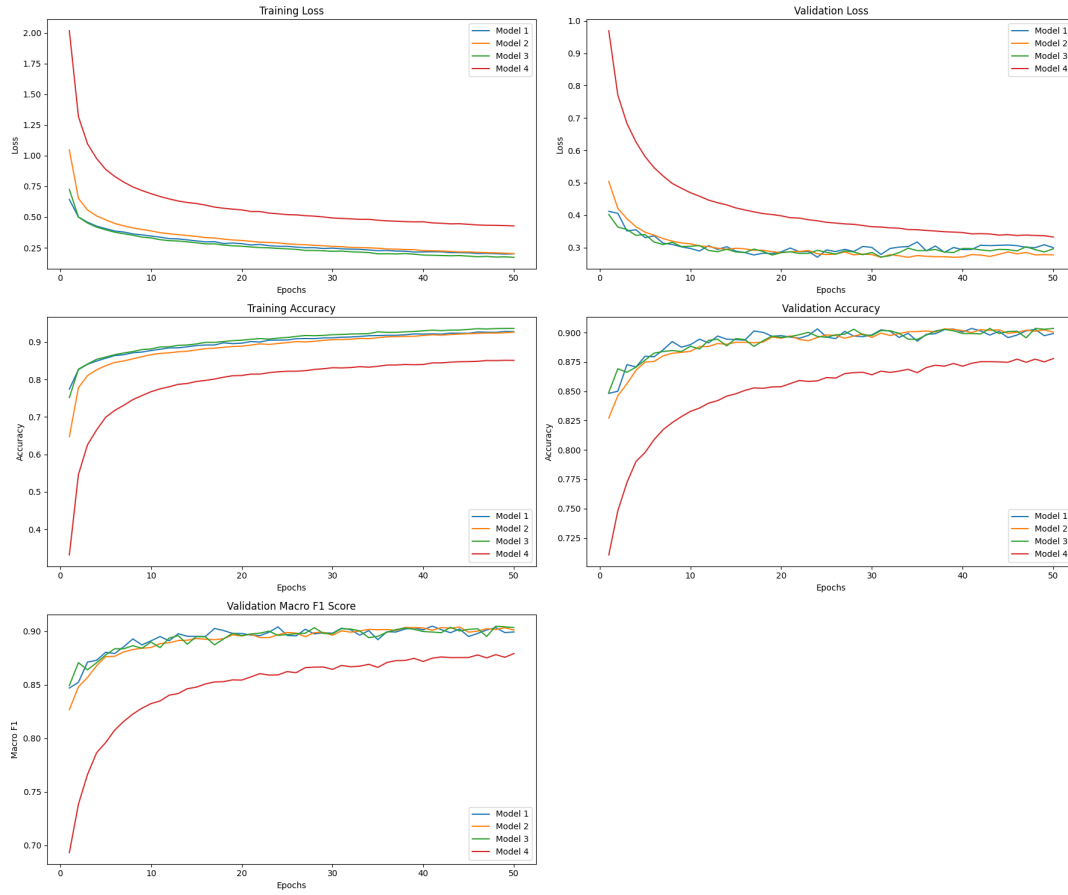


Figure 9: Comparison of Training and Validation History for All Four Models

## 3 Training the Final Model

The best-performing model from the validation phase was selected as the final Model.

### Configuration:

```
config_model_best = {  
    'learning_rate': 0.005,  
    'num_epochs': 50,  
    'batch_size': 32,  
    'hidden_layers': [  
        {'size': 512, 'bn': True, 'dropout': 0.3, 'relu': True},  
        {'size': 256, 'bn': False, 'dropout': 0.4, 'relu': True},  
        {'size': 128, 'bn': True, 'dropout': 0.5, 'relu': True}  
    ]  
}
```

### Training Performance:

- Train Accuracy: 0.9283
- Train Loss: 0.1972

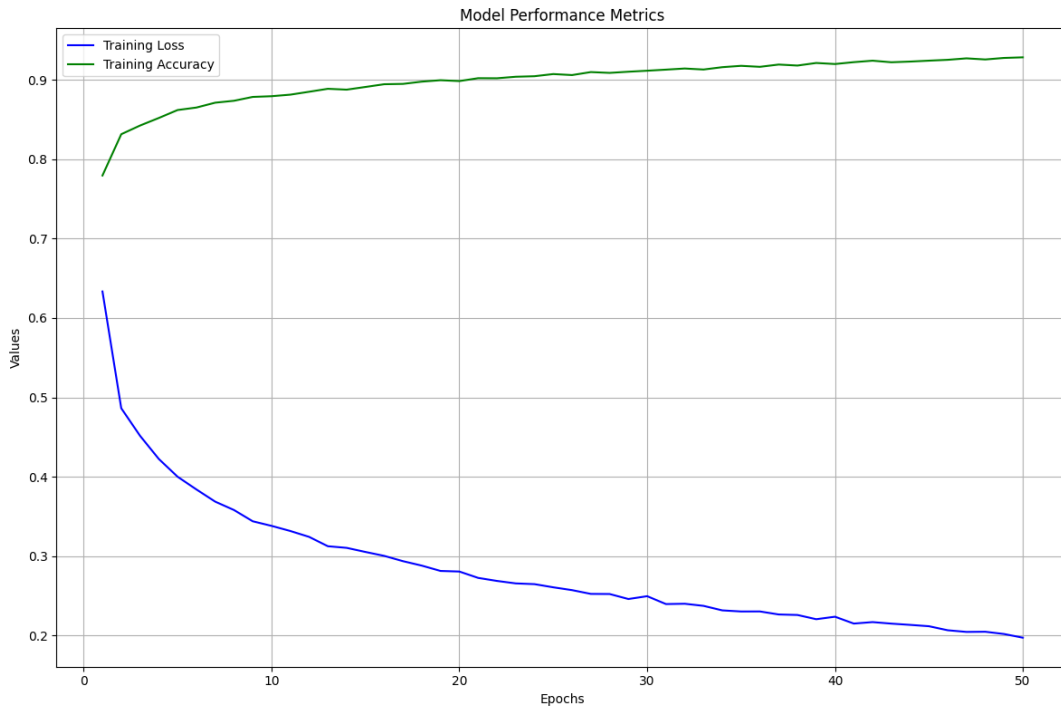


Figure 10: Training History of Best Model

## 4 Independent Testing Performance

Independent Testing on Final Model. **Performance:**

- Test Accuracy: 0.9066
- Test Loss: 0.3026
- Test Macro F1 Score: 0.9067

**Confusion Matrix:**

Table 8: Confusion Matrix

	0	1	2	3	4	5	6	7	8	9
0	828	0	11	21	4	1	128	0	7	0
1	3	985	1	7	0	0	2	0	2	0
2	15	1	854	7	57	0	65	0	1	0
3	14	7	7	921	31	0	17	0	3	0
4	0	0	78	25	834	1	60	0	2	0
5	0	0	0	1	0	982	0	9	1	7
6	94	0	59	26	56	0	760	0	5	0
7	0	0	0	0	0	19	0	957	0	24
8	5	0	5	3	1	3	2	3	978	0
9	0	0	0	0	0	8	1	24	0	967